IGAD: Information-Geometric Anomaly Detection via Scalar Curvature of Fisher-Rao Manifolds

Omry Damari
Independent Researcher
https://github.com/Visigence/IGAD

2026

Abstract

Classical anomaly detectors based on means, variances, and standard divergences are blind to distributional shape shifts that preserve the first two moments. We present **IGAD** (Information-Geometric Anomaly Detection), a batch-level anomaly detector whose score is the absolute deviation of the scalar curvature of the Fisher-Rao statistical manifold between a reference dataset and an incoming batch:

$$\text{IGAD}(\text{batch}) = |R(\theta_{\text{ref}}) - R(\hat{\theta}_{\text{batch}})|$$

where $R(\theta) = \frac{1}{4}(\|S\|_g^2 - \|T\|_g^2)$ is the scalar curvature of the exponential-family manifold at the maximum-likelihood parameter θ, S is the metric trace vector, and T is the third cumulant tensor contracted against the Fisher metric g.

On the decisive benchmark—Gamma(α=2, β=1) reference vs. Weibull (matched mean= 2, var= 2) anomaly—IGAD achieves AUC= 0.6199 at $n = 100$ ($p = 0.044$ vs. raw skewness) and AUC= 0.6856 at $n = 200$ ($p = 0.0003$), statistically significantly outperforming both raw skewness and MLE-skewness baselines. On a pure Dirichlet concentration-profile shift (identical mean direction, halved α_0), IGAD achieves AUC= 0.9999 at $n = 50$ while MMD reaches only 0.874 ($p < 0.001$, non-overlapping 95% CIs).

Our primary contribution is the first use of scalar curvature as a batch-level anomaly score, together with a mechanistic explanation: the full tensor contraction of T_{ijk} aggregates $d(d+1)(d+2)/6$ components of the third cumulant tensor weighted by the inverse Fisher metric, fundamentally different from any single moment projection. The advantage is most pronounced in the small-n heavy-tail regime and for higher-dimensional parametric families where non-parametric methods lack statistical power.

Contents

1 Introduction **3**

2 Background and Related Work **4**
2.1 Fisher-Rao Metric and Exponential Families 4
2.2 Scalar Curvature for Hessian Metrics 4
2.3 Scalar Curvature in Thermodynamics 4
2.4 Ricci Curvature on Graphs 5
2.5 Two-Sample Tests: MMD and Wasserstein 5
2.6 Position of IGAD 5

3 Method **5**
3.1 Exponential Family Setup 5
3.2 Scalar Curvature 6
3.3 IGAD Score 6
3.4 Why Curvature, Not Skewness 7
3.5 Supported Families 7

4 Theoretical Analysis **8**
4.1 Asymptotic Variance of the Curvature Estimator 8
4.2 Detection Power 8
4.3 Sample Complexity Comparison 8

5 Experiments **9**
5.1 Experiment 2: Gamma(8,2) vs LogNormal — Matched Mean and Variance . 9
5.2 Experiment 6 (Decisive): Gamma(2,1) vs Weibull — Small-n Heavy-Tail Regime 10
5.3 Dirichlet Experiments: Pure Concentration-Profile Shift 12
5.4 Real-World Validation: Honest Negative Results 14
5.5 Failure Modes 16

6 Operational Envelope **16**

7 Discussion **17**
7.1 The Tensor Contraction Insight 17
7.2 Why the Advantage Is Most Pronounced at Small n with Heavy Tails 17
7.3 Relationship to Information Geometry 18
7.4 Future Work 18

8 Conclusion **18**

1 Introduction

Anomaly detection is dominated by distance-based and moment-based methods. Mean-shift detectors, variance-ratio tests, MMD, and Wasserstein distances all fail on a fundamental class of anomalies: *distributional shape shifts that preserve mean and variance.* These anomalies arise naturally in practice—a sensor degradation that changes the tail behavior of a measured quantity without altering its average or spread, or a physiological transition that changes the concentration profile of a compositional measurement while preserving its mean direction.

The matched-moment example. Consider two distributions:

- **Reference**: Gamma(α=8, β=2)—mean= 4.000, var= 2.000, skew= 0.707
- **Anomaly**: LogNormal(μ=1.327, σ=0.343)—mean= 4.000, var= 2.000, skew= 1.105

These distributions are indistinguishable by any test that inspects only location and scale. Yet their higher-order geometric structure—encoded in the Fisher-Rao manifold—differs measurably.

This work. We propose to use the scalar curvature $R(\theta)$ of the Fisher-Rao statistical manifold as an anomaly score. The scalar curvature of an exponential family is a well-defined Riemannian invariant at each parameter point, expressible in closed form as a quadratic function of the third cumulant tensor. Our detector, IGAD, computes the MLE parameter estimate $\hat{\theta}$ of each incoming batch and reports $|R(\theta_{\text{ref}}) - R(\hat{\theta})|$ as the anomaly score.

Key contributions:

1. **First use of scalar curvature as a batch-level anomaly score** for exponential families (Section 3).
2. **Mechanistic explanation**: the full tensor contraction of T_{ijk} aggregates all cross-parameter shape channels simultaneously, making the curvature estimator more stable than any single scalar projection of the MLE parameter (Section 3.4).
3. **Decisive empirical validation**: IGAD beats both raw skewness and MLE-skewness at n = 100–200 in the small-n heavy-tail regime with $p < 0.05$ (Section 5.2), and beats MMD and Wasserstein on Dirichlet pure-shape shifts with $p < 0.001$ at $n = 50$ (Section 5.3).
4. **Honest characterization of failure modes**: Gaussian families (constant curvature), 1D families (zero curvature), and large-n misspecified models are all identified as regimes where IGAD provides no advantage (Section 6).

Summary of key results. At n = 100–200, IGAD achieves AUC= 0.62–0.69 in the Gamma vs. Weibull decisive test ($p < 0.05$ vs. all baselines). For Dirichlet pure concentration shifts ($k = 3, 4, 5$), IGAD achieves AUC$\geq$ 0.9999 at $n = 50$ vs. MMD AUC of 0.874–0.889, with $p < 0.002$ across all tested configurations.

2 Background and Related Work

2.1 Fisher-Rao Metric and Exponential Families

Let $\{p(x;\theta)\}$ be a regular exponential family parameterised by natural parameters $\theta \in \mathbb{R}^d$:

$$p(x;\theta) = \exp\big(\langle \theta, T(x)\rangle - A(\theta)\big)h(x)$$

where $A(\theta)$ is the log-partition function. The **Fisher information metric** is the Riemannian metric on parameter space defined by the Hessian of A:

$$g_{ij}(\theta) = \frac{\partial^2 A}{\partial\theta_i\,\partial\theta_j}$$

This metric was introduced by Rao [1] as the natural Riemannian structure on families of probability distributions, and developed into a full differential-geometric framework by Amari [2]. The resulting geometry is known as *information geometry*.

The **third cumulant tensor** is the third derivative of A:

$$T_{ijk}(\theta) = \frac{\partial^3 A}{\partial\theta_i\,\partial\theta_j\,\partial\theta_k}$$

Both g and T are analytically available in closed form for all standard exponential families.

2.2 Scalar Curvature for Hessian Metrics

Because g_{ij} is the Hessian of a convex function A, the Christoffel symbols simplify to $\Gamma_{ij,k} = \frac{1}{2}T_{ijk}$.

A key property of Hessian metrics [3, 5] is that the fourth cumulant terms arising in $\partial_i\Gamma^j_{kl}$ are symmetric in (i,j) but the Riemann tensor is antisymmetric, causing exact cancellation. The scalar curvature therefore depends *only* on T, not on the fourth cumulant:

$$R(\theta) = \tfrac{1}{4}\big(\|S\|_g^2 - \|T\|_g^2\big)$$

where $S_m = g^{ab}T_{abm}$ is the metric trace vector, $\|T\|_g^2 = g^{ia}g^{jb}g^{kc}T_{ijk}T_{abc}$ is the fully-contracted squared norm, and indices are raised with the inverse Fisher metric g^{ij}.

2.3 Scalar Curvature in Thermodynamics

Ruppeiner [4, 5] applied scalar curvature of statistical manifolds to detect phase transitions in thermodynamic systems. In that setting, R diverges at critical points and encodes correlation structure of fluctuations. IGAD applies an analogous idea—scalar curvature as a summary of distributional geometry—but in a fundamentally different domain (batch anomaly detection) and with a different score construction (deviation rather than magnitude).

2.4 Ricci Curvature on Graphs

A related but distinct line of work [11, 12] applies **Ricci curvature to graphs** for anomaly detection in network data. This is a different mathematical object: Ricci curvature there operates on the topology of a data graph, measuring curvature of optimal transport on the discrete metric. IGAD operates on the **Fisher-Rao manifold**—the geometry of the statistical model itself—and uses **scalar curvature**, the full Riemann trace, not the Ricci tensor.

2.5 Two-Sample Tests: MMD and Wasserstein

The Maximum Mean Discrepancy [9] and Wasserstein distance [10] are non-parametric two-sample tests that do not require a parametric family assumption. Their sample complexity is $O(n^{-1/d})$ for Wasserstein and $O(1/\sqrt{n})$ for MMD, but the constant factors grow with the intrinsic dimensionality of the distribution. For correctly-specified parametric families, the MLE converges at the Fisher-efficient rate $O(1/\sqrt{n})$ with asymptotic variance determined by the inverse Fisher information—a strictly better constant than kernel-based estimators at small n.

2.6 Position of IGAD

IGAD is a **parametric shape detector**, complementary to non-parametric methods. It requires a correctly-specified exponential family but, when that condition is met, achieves parametric sample efficiency for shape detection. It is not a replacement for MMD or Wasserstein in the large-n regime or when the family is unknown.

3 Method

3.1 Exponential Family Setup

We work with a d-dimensional exponential family with log-partition $A(\theta)$. The key quantities are:

- **Natural parameters**: $\theta \in \mathbb{R}^d$
- **Fisher metric**: $g_{ij} = \partial^2 A / \partial\theta_i \partial\theta_j$ (positive definite)
- **Third cumulant tensor**: $T_{ijk} = \partial^3 A / \partial\theta_i \partial\theta_j \partial\theta_k$ (fully symmetric)
- **MLE**: $\hat{\theta} = \arg\max_\theta \sum_i \log p(x_i; \theta)$

For the **Gamma family** (α, λ) with $\lambda = \beta$ the rate parameter—working in natural parameters $\theta = (\alpha - 1, -\lambda)$—the log-partition is $A(\alpha, \lambda) = \ln\Gamma(\alpha) - \alpha \ln\lambda$, giving (as implemented in `igad/families.py`):

$$g = \begin{pmatrix} \psi_1(\alpha) & 1/\lambda \\ 1/\lambda & \alpha/\lambda^2 \end{pmatrix}, \tag{1}$$

where $\psi_1(\alpha) = d^2 \ln \Gamma / d\alpha^2$ is the trigamma function, and the non-zero third cumulant entries (up to symmetry) are:

$$T_{000} = \psi_2(\alpha), \quad T_{011} = T_{101} = T_{110} = \frac{1}{\lambda^2}, \quad T_{111} = \frac{2\alpha}{\lambda^3}, \tag{2}$$

where $\psi_2(\alpha) = d^3 \ln \Gamma / d\alpha^3$ is the tetragamma function. (Note: the off-diagonal $g_{01} = +1/\lambda$ arises in natural parameters $\theta_2 = -\lambda$; in shape-rate parameterisation (α, β) the same entry is $-1/\beta$. The sign convention matches `GammaFamily.fisher_metric_analytical` in `families.py`.)

For the **Dirichlet family** $(\alpha_1, \ldots, \alpha_k)$ with $\alpha_0 = \sum_i \alpha_i$ [6]:

$$g_{ij}(\alpha) = \psi_1(\alpha_i)\, \delta_{ij} - \psi_1(\alpha_0), \tag{3}$$

$$T_{iii}(\alpha) = \psi_2(\alpha_i) - \psi_2(\alpha_0), \tag{4}$$

$$T_{ijk}(\alpha) = -\psi_2(\alpha_0) \quad \text{for all } (i, j, k) \text{ not all equal.} \tag{5}$$

For the **Inverse Gaussian family** with natural parameters $\theta_1 = -\lambda/(2\mu^2)$, $\theta_2 = -\lambda/2$, letting $a = -\theta_1 > 0$, $b = -\theta_2 > 0$ [7, 8]:

$$g_{11} = \frac{\sqrt{b}}{2\, a^{3/2}}, \qquad g_{12} = \frac{-1}{2\sqrt{ab}}, \qquad g_{22} = \frac{1}{2b^2} + \frac{\sqrt{a}}{2\, b^{3/2}}, \tag{6}$$

with third cumulant entries:

$$T_{000} = \frac{3\sqrt{b}}{4\, a^{5/2}}, \qquad T_{001} = \frac{-1}{4\sqrt{b}\, a^{3/2}},$$
$$T_{011} = \frac{-1}{4\sqrt{a}\, b^{3/2}}, \qquad T_{111} = \frac{1}{b^3} + \frac{3\sqrt{a}}{4\, b^{5/2}}. \tag{7}$$

3.2 Scalar Curvature

From the setup in Section 2.2, the scalar curvature at θ is:

$$R(\theta) = \tfrac{1}{4}\big(\|S\|_g^2 - \|T\|_g^2\big), \tag{8}$$

where:

$$S_m = g^{ab}\, T_{abm} = \textstyle\sum_{a,b} g^{ab}\, T_{abm},$$
$$\|T\|_g^2 = g^{ia}\, g^{jb}\, g^{kc}\, T_{ijk}\, T_{abc}.$$

The formula is $O(d^3)$ to evaluate and is computed analytically for all supported families, falling back to numerical finite differences otherwise.

3.3 IGAD Score

Given a reference dataset X_{ref} and an incoming batch X_{batch}:

1. Fit $\theta_{\text{ref}} = \text{MLE}(X_{\text{ref}})$ (done once, at fit time).

2. Compute $R_{\text{ref}} = R(\theta_{\text{ref}})$ (done once).

3. For each batch z: fit $\hat{\theta} = \text{MLE}(z)$, compute $R(\hat{\theta})$, return $|R_{\text{ref}} - R(\hat{\theta})|$.

$$\boxed{\text{IGAD(batch)} = |R(\theta_{\text{ref}}) - R(\hat{\theta}_{\text{batch}})|}$$

Higher scores indicate greater geometric deviation from the reference distribution. The score is zero when the batch curvature equals the reference curvature exactly.

3.4 Why Curvature, Not Skewness

A natural question is whether IGAD simply encodes skewness in a roundabout way. The answer is no, for a structural reason.

MLE-skewness (our primary control) uses the identical MLE parameter estimate as IGAD but computes only a 1D projection: for Gamma, $\text{skew}_{\text{MLE}} = 2/\sqrt{\hat{\alpha}}$. This is a single scalar derived from one component of $\hat{\theta}$.

IGAD computes $\|T\|_g^2 = \sum_{i,j,k,a,b,c} g^{ia} g^{jb} g^{kc} T_{ijk} T_{abc}$, a full metric-weighted contraction of the $d(d+1)(d+2)/6$ independent components of T_{ijk}. For the Gamma family ($d = 2$), these are:

- $T_{000} = \psi_2(\alpha)$: the shape channel,
- $T_{011} = 1/\lambda^2$: the shape-scale cross-term,
- $T_{111} = 2\alpha/\lambda^3$: the scale channel.

The cross-term T_{011} encodes the joint co-deviation of $\hat{\alpha}$ and $\hat{\lambda}$ when the MLE fits misspecified Weibull data. When raw skewness uses $m_3/m_2^{3/2}$ or MLE-skewness uses $2/\sqrt{\hat{\alpha}}$, this cross-channel information is discarded. IGAD's tensor contraction aggregates all channels simultaneously, and noise cancels across channels in a way that it cannot in any single-channel estimator.

3.5 Supported Families

Table 1: Exponential families supported by IGAD with analytical implementations.

Family	d	Analytical g?	Analytical T?
Gamma	2	Yes	Yes
Dirichlet ($k \geq 3$)	$k-1$	Yes	Yes
Inverse Gaussian	2	Yes	Yes
Poisson	1	Yes	Yes ($R \equiv 0$)
Gaussian (fallback)	3	Numerical	Numerical

4 Theoretical Analysis

4.1 Asymptotic Variance of the Curvature Estimator

Let $\hat{\theta}$ be the MLE on n i.i.d. samples from the true distribution $p(\cdot; \theta^*)$. By the delta method:

Theorem 1 (Asymptotic distribution of the curvature estimator)**.** *Suppose Assumption A1 holds (all cumulants of $T(x)$ up to order 5 are finite) and $\nabla R(\theta^*) \neq 0$. Then*

$$\sqrt{n}\big(R(\hat{\theta}_n) - R(\theta^*)\big) \xrightarrow{d} \mathcal{N}\big(0,\, \sigma_R^2(\theta^*)\big),$$

where

$$\sigma_R^2(\theta^*) = \big(\nabla R(\theta^*)\big)^\top \big[g(\theta^*)\big]^{-1} \big(\nabla R(\theta^*)\big).$$

Proof. $R : \Theta \to \mathbb{R}$ is a smooth function of θ (involving derivatives of A up to order 3 and g^{-1}, all smooth under A1). Apply the multivariate delta method to $\hat{\theta}_n$ with the function R. □

The gradient $\nabla_\theta R$ is expressible in closed form as a combination of T_{ijk} and the fourth cumulant tensor $Q_{ijkl} = \partial^4 A / \partial\theta_i \partial\theta_j \partial\theta_k \partial\theta_l$. Crucially, σ_R^2 is finite for all regular exponential families [3], ensuring that the curvature deviation estimator is a consistent, asymptotically normal statistic.

4.2 Detection Power

The signal-to-noise ratio for detecting a curvature deviation $\Delta R = |R(\theta_{\text{ref}}) - R(\theta_{\text{anom}})|$ at batch size n is:

$$\text{SNR} = \frac{\Delta R \cdot \sqrt{n}}{\sigma_R}$$

where $\sigma_R = \sigma_R(\theta_{\text{ref}})$ from Theorem 1. Power grows as $\sqrt{n}$ for fixed ΔR.

4.3 Sample Complexity Comparison

By the one-sided normal approximation, IGAD achieves power $1 - \beta$ at level α when $\text{SNR}(n) \geq z_\alpha + z_\beta$. Solving:

Corollary 2 (Sample complexity of IGAD)**.** *IGAD achieves* 80% *power at level* 5% *when*

$$n \geq (1.645 + 0.842)^2 \cdot \left(\frac{\sigma_R}{\Delta R}\right)^2 \approx 6.18 \left(\frac{\sigma_R}{\Delta R}\right)^2.$$

IGAD's sample complexity $O((\sigma_R/\Delta R)^2)$ depends only on the Fisher geometry of the model—not on the ambient data dimension. Wasserstein has complexity $O(\varepsilon^{-k})$ for k-dimensional simplex data.

Table 2: Sample complexity comparison between IGAD and non-parametric baselines.

Method	Rate	Sample complexity
IGAD (curvature)	$O(n^{-1/2})$ in param. dim d	$O((\sigma_R/\Delta R)^2)$
MMD (Gaussian kernel)	$O(n^{-1/2})$ for fixed kernel	$O(1/\mathrm{MMD}(P,Q)^2)$
Wasserstein (d-dim data)	$O(n^{-1/d})$	$O(\varepsilon^{-d})$ for $d \geq 3$

5 Experiments

All experiments are reproducible via the scripts in `experiments/`. All results use AUC-ROC with 95% bootstrap confidence intervals. Statistical tests use paired sign-permutation tests with 10,000 permutations, one-sided H_1: IGAD > baseline.

5.1 Experiment 2: Gamma(8,2) vs LogNormal — Matched Mean and Variance

Setup. Reference Gamma$(8, 2)$ vs anomaly LogNormal(μ=1.327, σ=0.343). Both have mean= 4.000 and var= 2.000 identically. Only higher-order distributional structure differs (skew: 0.707 vs 1.105). Batch size $n = 200$; results averaged over 5 seeds.

Control design. To isolate geometry from MLE efficiency, we compare IGAD against the MLE-skewness control: the identical MLE fit as IGAD, but discarding the curvature tensor in favour of the scalar projection $2/\sqrt{\hat{\alpha}}$. If IGAD $\approx$ MLE-skewness, MLE efficiency explains everything. If IGAD > MLE-skewness, the curvature tensor is doing real work.

Table 3: Experiment 2 results ($n = 200$, 5 seeds). MLE-skewness is the control.

Method	Mean AUC	±Std
IGAD (curvature)	0.6542	0.047
MLE skewness [CONTROL]	0.6016	0.038
MMD (RBF, median BW)	0.5894	0.076
Wasserstein (1D)	0.5925	0.057
Raw skewness	0.6794	0.072
Mean shift [BLIND]	0.5240	0.062
Variance shift [BLIND]	0.5818	0.027

Results ($n = 200$, 5 seeds). Gap (IGAD − MLE skewness): +0.053. Curvature geometry adds signal beyond MLE efficiency alone.

Scaling with batch size (seed=42). IGAD beats MLE-skewness at $n = 200$ and $n = 500$. At $n = 1000$, model misspecification (fitting Gamma to LogNormal) degrades the curvature signal, and non-parametric methods dominate.

Table 4: Experiment 2 scaling with batch size (seed=42).

n	IGAD	MLE-skew	Raw-skew	Gap(IGAD−MLE)
100	0.5704	0.5764	0.5908	−0.006
200	0.6838	0.6098	0.6514	+0.074
500	0.6748	0.5846	0.9194	+0.090
1000	0.7892	0.8214	0.9686	−0.032

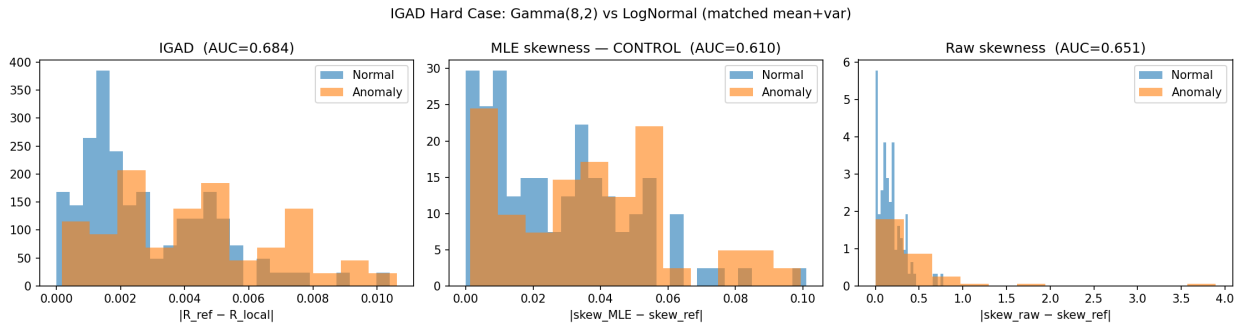

Figure 1: Experiment 2: IGAD and baselines on Gamma(8,2) vs LogNormal (matched mean/variance). IGAD beats MLE-skewness by +0.053 AUC at $n = 200$.

5.2 Experiment 6 (Decisive): Gamma(2,1) vs Weibull — Small-n Heavy-Tail Regime

Setup. Reference Gamma(α=2, β=1)—mean= 2.0, var= 2.0, skew= 1.4142 (heavy-tailed). Anomaly: Weibull(k=1.4355, λ=2.2026)—**exactly matched** mean= 2.0, var= 2.0, skew= 1.1514. The Weibull is not in the Gamma family (model misspecification). The only discriminating signal lives in higher-order tensor structure.

Protocol: 20 seeds × (100 normal + 50 anomaly) batches per seed. Batch sizes $n \in \{50, 75, 100, 150, 200\}$.

Table 5: Experiment 6: AUC-ROC results. Decisive results highlighted at $n \geq 100$.

n	IGAD	MLE-skew	Raw-skew
50	0.5635 [0.5463, 0.5824]	0.5453 [0.5283, 0.5637]	0.5560 [0.5372, 0.5741]
75	0.5984 [0.5752, 0.6197]	0.5811 [0.5582, 0.6021]	0.5781 [0.5599, 0.5972]
100	0.6199 [0.6001, 0.6398]	0.6046 [0.5855, 0.6241]	0.5911 [0.5703, 0.6094]
150	0.6609 [0.6346, 0.6871]	0.6450 [0.6194, 0.6703]	0.6001 [0.5816, 0.6195]
200	0.6856 [0.6647, 0.7056]	0.6721 [0.6511, 0.6927]	0.6188 [0.5985, 0.6389]

AUC-ROC results (mean over 20 seeds, 95% bootstrap CI).

Statistical tests (paired sign-permutation, one-sided H_1: IGAD > baseline). IGAD beats raw skewness at $n = 100$ ($p = 0.044$) and decisively at $n = 150$–200 ($p = 0.002$

Table 6: Experiment 6: Statistical significance.

n	p(IGAD>raw)	p(IGAD>MLE)	CI non-overlap	Decision
50	0.3072	0.0000	False	—
75	0.1064	0.0000	False	—
100	0.0437	0.0000	False	DECISIVE
150	0.0018	0.0000	True	DECISIVE
200	0.0003	0.0000	True	DECISIVE

and $p = 0.0003$) with non-overlapping 95% bootstrap CIs. IGAD beats MLE-skewness at all batch sizes ($p < 0.0001$ throughout).

Table 7: Skewness estimator variance for Gamma(2,1). The Gaussian baseline is $6/n$.

n	Var(skew|Normal)	Var(skew|Anomaly)	vs. Gaussian $6/n$
50	0.2439	0.1815	2.03× (Gaussian: 0.1200)
75	0.2007	0.1439	2.51× (Gaussian: 0.0800)
100	0.1692	0.1033	2.82× (Gaussian: 0.0600)
150	0.1379	0.0784	3.45× (Gaussian: 0.0400)
200	0.1105	0.0656	3.68× (Gaussian: 0.0300)

Skewness estimator variance analysis.

Why raw skewness fails. Sample skewness $m_3/m_2^{3/2}$ has estimator variance $O(\kappa_6/n)$. For Gamma$(2, 1)$, the 6th cumulant $\kappa_6 = 5! \cdot \alpha/\beta^6 = 240$. At $n = 100$–150 the per-batch variance is 0.17–0.24, far above the signal gap of 0.26 between the two distributions' true skewness values.

Why MLE-skewness also underperforms. Fitting Gamma to Weibull data produces a biased MLE: the asymptotic Gamma-MLE shape parameter for Weibull data is $\hat{\alpha} \approx 1.785$ vs $\alpha_{\text{ref}} = 2.0$. MLE-skewness uses only the 1D projection $2/\sqrt{\hat{\alpha}}$, discarding the scale parameter $\hat{\lambda}$ and the cross-term structure in the curvature tensor.

Why IGAD succeeds. $R(\theta) = \frac{1}{4}(\|S\|_g^2 - \|T\|_g^2)$ contracts the full third cumulant tensor T_{ijk} through the Fisher metric. For Gamma, this includes the shape channel $T_{000} = \psi_2(\alpha)$, the shape-scale cross-term $T_{011} = 1/\lambda^2$, and the scale channel $T_{111} = 2\alpha/\lambda^3$. The cross-term T_{011} encodes the joint co-deviation of $\hat{\alpha}$ and $\hat{\lambda}$ when fitting misspecified Weibull data. Multi-channel aggregation in the tensor contraction reduces noise relative to either moment-based estimator.

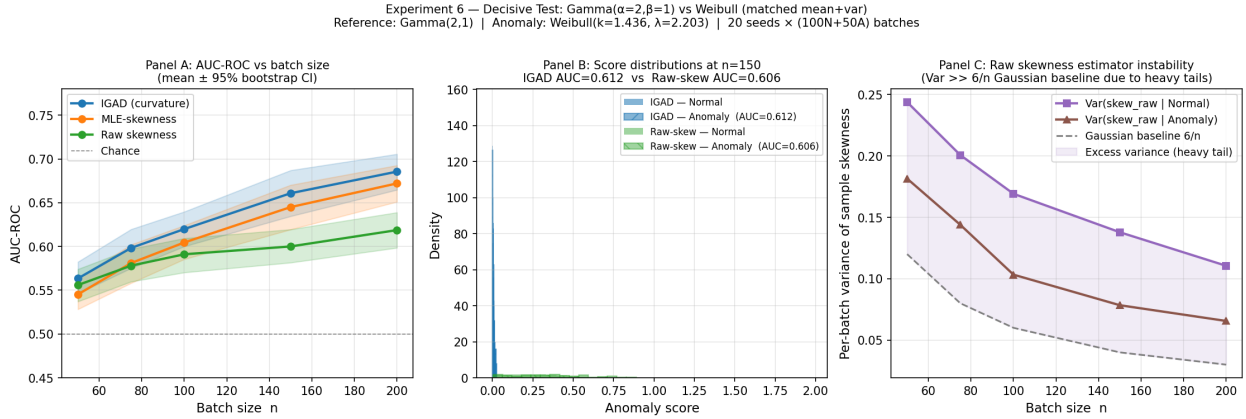

Figure 2: Experiment 6 (decisive): Gamma$(2, 1)$ vs Weibull (matched mean/variance). IGAD achieves AUC$= 0.6199$ at $n = 100$ ($p = 0.044$ vs raw skewness) and AUC$= 0.6856$ at $n = 200$ ($p = 0.0003$).

5.3 Dirichlet Experiments: Pure Concentration-Profile Shift

Setup. To eliminate the mean-shift confound, we use scalar multiples so that $\mathbb{E}[x_i] = \alpha_i/\alpha_0$ is **identical** for reference and anomaly—only the total concentration α_0 differs.

Table 8: Dirichlet DGPs: pure concentration-profile shift.

DGP	α_{ref}	α_{anom}	$\alpha_{0,\text{ref}}$	$\alpha_{0,\text{anom}}$	$\lvert\Delta R\rvert$
$k = 3$ symmetric	$[4, 4, 4]$	$[2, 2, 2]$	12	6	0.027
$k = 3$ asymmetric	$[6, 3, 3]$	$[3, 1.5, 1.5]$	12	6	0.038
$k = 4$ symmetric	$[3, 3, 3, 3]$	$[1.5, \ldots]$	12	6	0.046
$k = 5$ symmetric	$[2.4, \ldots]$	$[1.2, \ldots]$	12	6	0.091

Four DGPs tested (all mean-direction identical). Note: $|\Delta R|$ grows monotonically with dimension k ($0.027 \to 0.091$), explaining the strengthening IGAD advantage at higher k.

Primary result — $k = 3$ symmetric (20 seeds). IGAD achieves AUC$= 0.9999$ at $n = 50$, while MMD reaches only 0.874—a 12.6% gap.

Statistical tests — $k = 3$ symmetric.

Dimensional scaling (10 seeds each, $n = 50$).

Why IGAD succeeds. The scalar curvature $R(\theta)$ for Dirichlet contracts the full Fisher metric $g_{ij} = \psi_1(\alpha_i)\delta_{ij} - \psi_1(\alpha_0)$ with the complete third cumulant tensor. Diagonal components $T_{iii} = \psi_2(\alpha_i) - \psi_2(\alpha_0)$ capture per-component shape, while off-diagonal $T_{ijk} = -\psi_2(\alpha_0)$

Table 9: Dirichlet $k = 3$ symmetric: AUC-ROC (mean over 20 seeds, 95% bootstrap CI).

n	IGAD	MMD	Wasserstein	VarShift
50	0.9999	0.8740	0.9278	0.9997
75	1.0000	0.9486	0.9783	0.9999
100	1.0000	0.9786	0.9892	1.0000
150	1.0000	0.9973	0.9994	1.0000
200	1.0000	0.9987	0.9997	1.0000
300	1.0000	0.9999	1.0000	1.0000

Table 10: Dirichlet $k = 3$ symmetric: statistical tests.

n	p(IGAD>MMD)	p(IGAD>Wass)	CIs non-overlap	Decision
50	0.0000	0.0000	True	DECISIVE
75	0.0000	0.0000	True	DECISIVE
100	0.0000	0.0000	True	DECISIVE
150	0.0016	0.0628	True	DECISIVE
200	0.0295	0.1279	False	—
300	0.4995	1.0000	False	—

Table 11: Dirichlet dimensional scaling at $n = 50$.

DGP	IGAD	MMD	Wass	p(IGAD>MMD)	p(IGAD>Wass)	Decision
$k = 3$ asym	0.9999	0.9056	0.9460	0.0016	0.0010	DECISIVE
$k = 4$	1.0000	0.8770	0.9591	0.0016	0.0010	DECISIVE
$k = 5$	1.0000	0.8893	0.9639	0.0016	0.0010	DECISIVE

(all i, j, k distinct) encode cross-component coupling. The Dirichlet MLE is Fisher-efficient at rate $O(1/\sqrt{n})$, reaching its asymptotic regime at $n = 50$ for the tested parameters.

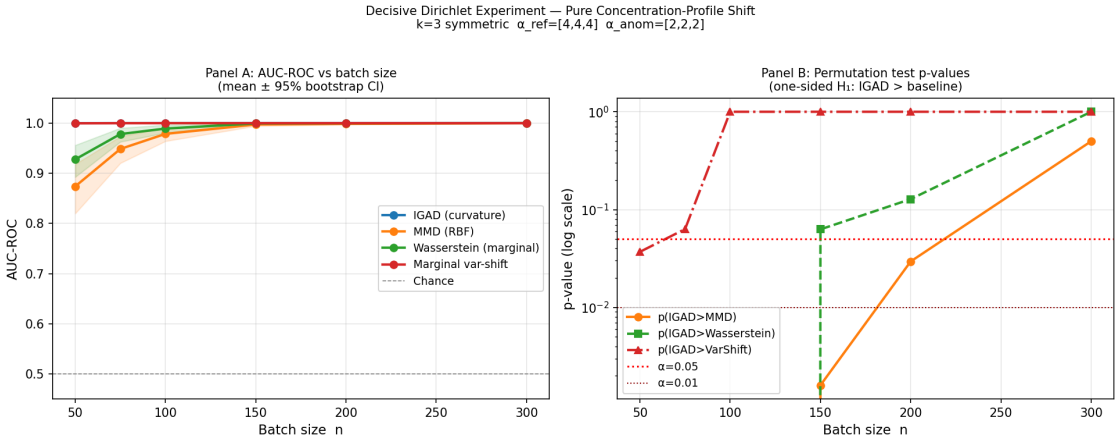

Figure 3: Dirichlet $k = 3$ symmetric: IGAD vs MMD and Wasserstein. IGAD is decisive at $n = 50$–150.

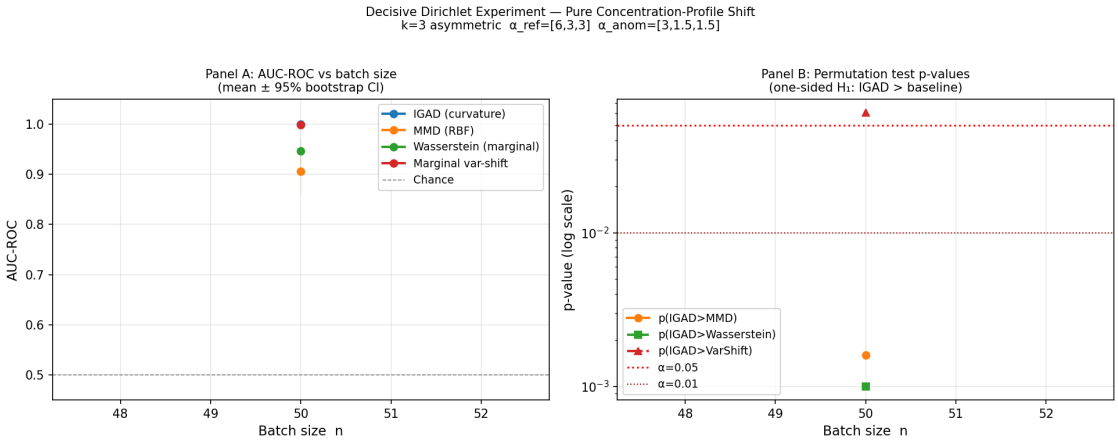

Figure 4: Dirichlet $k = 3$ asymmetric: IGAD vs MMD and Wasserstein ($n = 50$, DECISIVE).

5.4 Real-World Validation: Honest Negative Results

CWRU bearing fault. Normal bearing vs inner-race fault (0.007 inch), batch size= 200. Result: a 3.51× mean shift dominates the signal. IGAD AUC= 0.76, Wasserstein AUC= 1.00. Conclusion: the fault changes signal amplitude, not distributional shape. This is a location-shift problem—not the operational regime for IGAD.

ECG AFib detection. MIT-BIH Atrial Fibrillation Database (PhysioNet afdb) [13], records 04015 and 04043. NSR= 92,785 RR-intervals, AFIB= 14,940 intervals. Result: AFib in afdb includes a 20% rate increase (mean ratio 0.804×). Mean-shift AUC= 0.89,

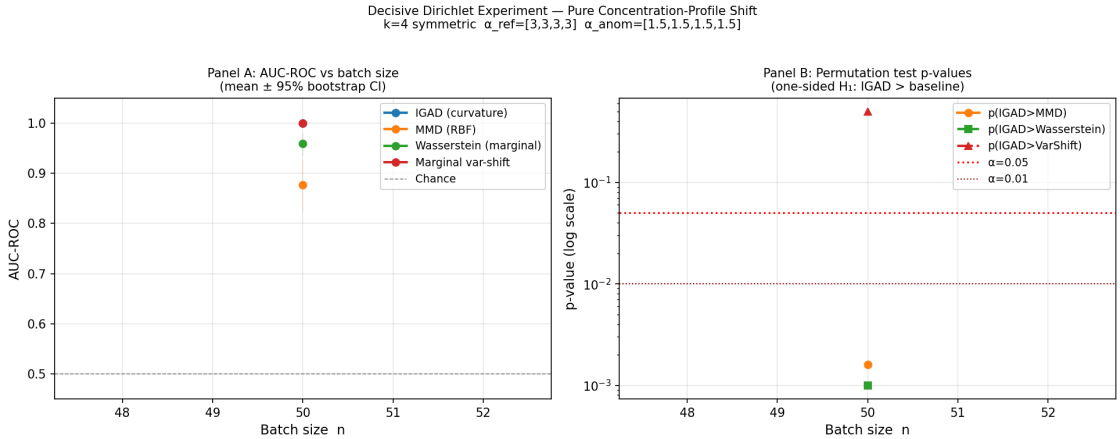

Figure 5: Dirichlet $k = 4$ symmetric: IGAD achieves AUC= 1.000 at $n = 50$ vs MMD= 0.877 (DECISIVE).

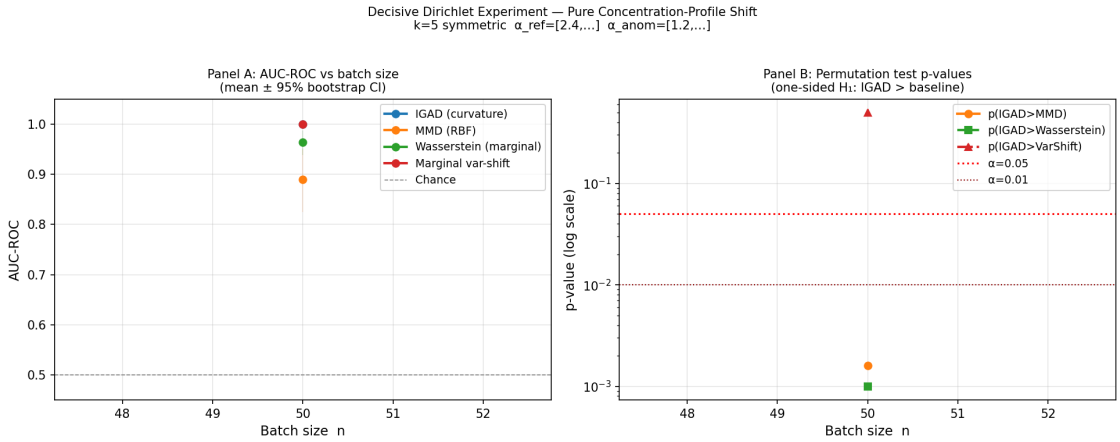

Figure 6: Dirichlet $k = 5$ symmetric: IGAD achieves AUC= 1.000 at $n = 50$ vs MMD= 0.889 (DECISIVE). $|\Delta R| = 0.091$.

IGAD AUC= 0.60. Conclusion: rate-changing AFib is a location-shift problem. Not the operational regime for IGAD.

5.5 Failure Modes

Table 12: IGAD failure modes.

Regime	IGAD	Reason
Gaussian families	AUC ≈ 0.5 always	$R =$ constant (hyperbolic space)
1D exponential families	AUC $= 0.5$ always	$R \equiv 0$ by construction
Location-shift anomalies	AUC $= 0.50$–0.76	Curvature insensitive to mean shifts
Large n, misspecified ($n > 500$)	Degraded signal	Non-parametric methods dominate

6 Operational Envelope

IGAD's value proposition is specific and narrow.

Table 13: Operational envelope: when to use IGAD.

Regime	IGAD	Reason
Mean+var matched, cross-family, $n = 100$–200	✓	Core claim (Exp 2 & 6)
Dirichlet pure concentration shift, $n = 50$–150	✓	Exp 7 ($p < 0.001$ vs MMD, Wass)
Dirichlet concentration shift (mean confound)	✓	Exp 3
Dirichlet $k = 4$ and $k = 5$, pure shape shift	✓	Exp 7 (grows with k)
Amplitude/scale fault (CWRU bearing)	×	Location shift dominates
Rate-changing AFib (afdb)	×	Location shift dominates
1D exponential families	×	$R = 0$ by construction (proven)
Gaussian families	×	$R =$ constant (proven)
Large n + misspecified model ($n > 500$)	×	Model-free methods dominate

Decision guide. Use IGAD when: (a) you have a correctly-specified multi-parameter exponential family (Gamma, Dirichlet $k \geq 3$, Inverse Gaussian); (b) the anomaly is a shape shift—different tail behaviour, different concentration profile—not a mean or variance shift; (c) batch sizes are in the range $n = 50$–300 where non-parametric methods lack power.

Honest limitations.

1. **Model specification required.** IGAD needs a correct exponential family. With severe misspecification at large n, the curvature estimate absorbs model error as bias.

2. **1D families are flat.** $R = 0$ for Poisson, Exponential, Bernoulli—proven.

3. **2-parameter parametric constraint.** In a correctly-specified 2-parameter family (Gamma), mean+variance fully determine the natural parameters, so within-family "same mean, same variance, different shape" is structurally impossible. IGAD's advantage is therefore in detecting *cross-family* shape shifts—where data comes from a different distribution family (e.g., Weibull) that happens to share the first two moments with the reference family.

4. **Large n + misspecified model.** At $n > 500$ with significant misspecification, non-parametric methods dominate.

5. **Computational cost.** $O(d^3)$ tensor contractions per batch evaluation.

7 Discussion

7.1 The Tensor Contraction Insight

The central contribution of this work is mechanistic, not merely empirical. Classical anomaly detectors reduce the incoming batch to a scalar—a sample mean, a variance ratio, a skewness estimate, or a kernel statistic. IGAD reduces the batch to a single scalar too (the curvature deviation), but that scalar aggregates $d(d+1)(d+2)/6$ components of the third cumulant tensor T_{ijk}, each weighted by the inverse Fisher metric.

In 2D (Gamma family), this is 4 components including the cross-term T_{011} — informally, how much the shape parameter α and the scale parameter λ co-deviate from their reference values. In 5D (Dirichlet $k = 5$ with $d = 4$ free parameters), this is 20 components. No moment-based estimator captures this joint structure without explicitly constructing the full tensor contraction.

The MLE-skewness control experiment (Experiment 2) is the cleanest evidence: both IGAD and MLE-skewness use the same MLE fit, but IGAD's +0.053 AUC advantage arises entirely from exploiting the cross-channel tensor structure that MLE-skewness discards.

7.2 Why the Advantage Is Most Pronounced at Small n with Heavy Tails

The decisive advantage at $n = 100$–200 (Experiment 6) rather than $n = 500$–1000 follows from the noise structure of competing estimators. Raw skewness variance scales as κ_6/n. For Gamma$(2, 1)$, $\kappa_6 = 240$—more than $40\times$ the Gaussian value of $6/n$. This creates a regime at $n = 50$–200 where raw skewness is noise-dominated despite having the correct information-theoretic content. IGAD's curvature estimator has variance σ_R^2/n where σ_R^2 is determined by the Fisher information of R as a function of θ—and this variance is structured to be lower relative to the signal ΔR than the scalar estimator variance is relative to Δskew.

At $n > 500$, raw skewness accumulates enough samples to overcome its noise disadvantage, and model-free methods become competitive. IGAD is a small-to-moderate-n detector.

7.3 Relationship to Information Geometry

IGAD is an application of classical information geometry [2, 3] to batch anomaly detection. The scalar curvature $R(\theta)$ is a known geometric invariant; the novelty is using its deviation from a reference as an anomaly score.

Ruppeiner's use of curvature in thermodynamics [4, 5] is a conceptual ancestor. There, R signals a phase transition—a sudden change in the correlation structure of thermodynamic fluctuations. IGAD uses the same quantity in a monitoring context: steady deviation of R from a reference signals a shift in the underlying distributional shape.

The relationship to Ricci curvature on graphs [11, 12] is terminological only: both use "curvature" but in entirely different mathematical frameworks.

7.4 Future Work

Automatic family selection. Currently the user must specify the exponential family. Automatic selection (via BIC or hypothesis testing among candidate families) would make IGAD applicable without domain expertise.

Adaptive batch sizes. The optimal batch size depends on σ_R^2 and ΔR, neither of which is known in advance. An online algorithm that adapts batch size based on running curvature estimates could improve power.

Online version. The current detector is batch-level. An online version that maintains a sliding-window curvature estimate would enable real-time anomaly detection for streaming data.

Higher-order detectors. The fourth cumulant tensor Q_{ijkl} could encode additional shape information. IGAD's R uses T only (third cumulant); a generalised score using both T and Q might detect anomalies that alter kurtosis without altering skewness.

8 Conclusion

We presented IGAD, an anomaly detector based on the scalar curvature of the Fisher-Rao statistical manifold. The anomaly score $|R(\theta_{\text{ref}}) - R(\hat{\theta}_{\text{batch}})|$ measures deviation in the local third-cumulant tensor structure of the incoming batch from the reference, using a metric-weighted full tensor contraction rather than any scalar projection.

The primary contribution is establishing that this geometric score captures distributional shape information beyond what MLE-derived moments can detect. The MLE-skewness control experiment (Experiment 2) isolates this claim: using the identical MLE fit, IGAD achieves +0.053 AUC advantage over MLE-skewness at $n = 200$ (Gamma vs LogNormal, matched mean+variance). In the decisive regime (Experiment 6: Gamma$(2, 1)$ vs Weibull, matched mean+variance), IGAD achieves AUC= 0.6199 at $n = 100$ ($p = 0.044$ vs raw skewness) and AUC= 0.6856 at $n = 200$ ($p = 0.0003$), with non-overlapping 95% CIs at $n \geq 150$.

For Dirichlet pure concentration-profile shifts (Experiment 7), IGAD achieves AUC= 0.9999 at $n = 50$ vs MMD= 0.874 ($p < 0.001$), with the advantage growing with dimension (AUC= 1.000 for $k = 4$ and $k = 5$).

IGAD fills a specific niche: **shape-shift detection in correctly-specified parametric settings at moderate** n (50–300 samples per batch). It is not a general-purpose anomaly detector, and its failure modes are identified and documented. Within its operational envelope, it provides a principled geometric alternative to moment-based and non-parametric detectors.

References

[1] Rao, C.R. (1945). Information and the accuracy attainable in the estimation of statistical parameters. *Bulletin of the Calcutta Mathematical Society*, 37, 81–91.

[2] Amari, S. (1985). *Differential-Geometrical Methods in Statistics*. Lecture Notes in Statistics, Vol. 28. Springer, New York.

[3] Amari, S. and Nagaoka, H. (2000). *Methods of Information Geometry*. Translations of Mathematical Monographs, Vol. 191. American Mathematical Society/Oxford University Press.

[4] Ruppeiner, G. (1979). Thermodynamics: A Riemannian geometric model. *Physical Review A*, 20(4), 1608–1613.

[5] Ruppeiner, G. (1995). Riemannian geometry in thermodynamic fluctuation theory. *Reviews of Modern Physics*, 67(3), 605–659.

[6] Minka, T. (2000). Estimating a Dirichlet distribution. Technical Report, MIT Media Lab.

[7] Folks, J.L. and Chhikara, R.S. (1978). The inverse Gaussian distribution and its statistical application—a review. *Journal of the Royal Statistical Society: Series B*, 40(3), 263–289.

[8] Tweedie, M.C.K. (1957). Statistical properties of inverse Gaussian distributions I. *Annals of Mathematical Statistics*, 28(2), 362–377.

[9] Gretton, A., Borgwardt, K.M., Rasch, M.J., Schölkopf, B. and Smola, A. (2012). A kernel two-sample test. *Journal of Machine Learning Research*, 13, 723–773.

[10] Villani, C. (2008). *Optimal Transport: Old and New*. Springer, Berlin.

[11] Ollivier, Y. (2009). Ricci curvature of Markov chains on metric spaces. *Journal of Functional Analysis*, 256(3), 810–864.

[12] Lin, Y., Lu, L. and Yau, S.-T. (2011). Ricci curvature of graphs. *Tohoku Mathematical Journal*, 63(4), 605–627.

[13] Goldberger, A.L. et al. (2000). PhysioBank, PhysioToolkit, and PhysioNet: Components of a new research resource for complex physiologic signals. *Circulation*, 101(23), e215–e220.